

Thot Toolkit for Statistical Machine Translation

User Manual

Daniel Ortiz Martínez
dortiz@prhlt.upv.es
Universitat Politècnica de València

April 2015

CONTENTS

1	Introduction	1
1.1	Statistical Foundations	1
1.1.1	Statistical Machine Translation	1
1.1.2	Computer-Aided Translation	2
1.2	Toolkit Features	4
1.3	Distribution Details	5
1.4	Relation with Existing Software	5
1.5	Current Status	5
1.6	Documentation and Support	6
1.7	Citation	7
1.8	Sponsors	8
2	Installation	9
2.1	Basic Installation Procedure	9
2.2	Alternative Installation Options	10
2.3	Installation Including the CasMaCat Workbench	11
3	User Guide	13
3.1	Toolkit Overview	13
3.2	Training and Tuning Tools	14
3.2.1	Language Model Training	14
3.2.2	Translation Model Training	15
3.2.3	Basic Configuration File Generation	16
3.2.4	Parameter Tuning	16
3.2.5	Phrase Model Filtering	17
3.3	Search Tools	18
3.3.1	Fully Automatic Translation	18
3.3.2	Interactive Machine Translation	19
3.3.3	Phrase Alignment Generation	20
3.4	Additional Tools	20
3.4.1	Output Evaluation	20
3.5	Advanced Functionality	20
3.5.1	Online Learning	20
3.6	General Sample Uses	21
3.6.1	Training, Tuning and Translating	21
3.6.2	Online Learning	22
3.7	Troubleshooting	23

4	Background	25
	Bibliography	27

CHAPTER 1

INTRODUCTION

Thot is an open source toolkit for statistical machine translation. Originally, Thot incorporated tools to train phrase-based models. The new version of Thot now includes a state-of-the-art phrase-based translation decoder as well as tools to estimate all of the models involved in the translation process. In addition to this, Thot is also able to incrementally update its models in real time after presenting an individual sentence pair.

1.1 Statistical Foundations

In this section, the foundations of statistical machine translation and its use in computer-aided applications are very briefly described.

1.1.1 Statistical Machine Translation

The statistical approach to MT formalises the problem of generating translations under a statistical point of view. More formally, given a source sentence $f_1^J \equiv f_1 \dots f_j \dots f_J$ in the source language \mathcal{F} , we want to find its equivalent target sentence $e_1^I \equiv e_1 \dots e_i \dots e_I$ ^a in the target language \mathcal{Y} .

From the set of all possible sentences of the target language, we are interested in the one with the highest probability according to the following equation:

$$\hat{e}_1^I = \arg \max_{I, e_1^I} \{Pr(e_1^I | f_1^J)\} \quad (1.1)$$

where $Pr(e_1^I | f_1^J)$ represents the probability of translating f_1^J into e_1^I .

The early works on SMT were based on the use of *generative models*. A generative model is a full probability model of all statistical variables that are required to randomly generating observable data. Generative models decompose $Pr(e_1^I | f_1^J)$ applying the Bayes decision

^a f_j and e_i note the i 'th word and the j 'th word of the sentences f_1^J and e_1^I respectively.

rule. Taking into account that $Pr(f_1^J)$ does not depend on e_1^I we arrive to the following expression (Brown et al. 1993):

$$\hat{e}_1^I = \arg \max_{I, e_1^I} \{Pr(e_1^I) \cdot Pr(f_1^J | e_1^I)\} \quad (1.2)$$

where: $Pr(e_1^I)$ represents the probability of generating the target sentence, and $Pr(f_1^J | e_1^I)$ is the probability of generating e_1^I given f_1^J . Since the real probability distributions $Pr(e_1^I)$ and $Pr(f_1^J | e_1^I)$ are not known, they are approximated by means of parametric statistical models. Specifically, $Pr(e_1^I)$ is modelled by means of a *language model*, and $Pr(f_1^J | e_1^I)$ is modelled by means of a *translation model*. Current MT systems are based on the use of *phrase-based models* (Koehn et al. 2003) as translation models. Typically, the values of the parameters of such statistical models are obtained by means of the well-known *maximum-likelihood* estimation method.

More recently, alternative formalizations have been proposed. Such formalizations are based on the direct modelling of the posterior probability $Pr(e_1^I | f_1^J)$, replacing the generative models by *discriminative models*. Log-linear models use a set of feature functions $h_m(f_1^J, e_1^I)$ each one with its corresponding weight λ_m :

$$\hat{e}_1^I = \arg \max_{I, e_1^I} \left\{ \sum_{m=1}^M \lambda_m h_m(f_1^J, e_1^I) \right\} \quad (1.3)$$

The direct optimization of the posterior probability in the Bayes decision rule is referred to as *discriminative training* (Ney 1995). Since the features of regular SMT log-linear models are usually implemented by means of generative models, discriminative training is applied here only to estimate the weights involved in the log-linear combination. This process is typically carried out by means of the *minimum error rate training* (MERT) algorithm (Och 2003).

1.1.2 Computer-Aided Translation

Despite multiple and important advances obtained so far in the field of SMT, current MT systems are in many cases not able to produce ready-to-use texts. Indeed, MT systems usually require human intervention in order to achieve high-quality translations. Here we consider two different types of computer-aided translation applications based on SMT: post-editing and interactive machine translation.

Post-Editing the Output of Statistical Machine Translation

Post-editing (PE) involves making corrections and amendments to machine generated translations (see (TAUS-Project 2010) for a detailed study). PE is used when raw machine translation is not error-free, situation which is common for current MT technology. PE started being used in the late seventies mainly at some big institutions (such as the European Commission) and is currently gaining acceptance from translation companies. Currently, PE tends to be carried out via tools built for editing human generated translations, such as translation memories (some authors refer to this task as simply *editing*). In addition to this, new translation

memory tools and new versions of established ones offer translators the option to post-edit machine generated text for segments lacking any matches in the memories (Garcia 2011).

Since in the PE scenario, the user only edits the output of the MT system without further intervention from the system, there are no differences in the way in which the MT system is designed and implemented. Hence, the statistical framework for MT described above can be adopted without modifications in order to build the PE system.

Statistical Interactive Machine Translation

The interactive machine translation (IMT) framework constitutes an alternative to fully automatic MT systems in which the MT system and its user collaborate to generate correct translations. These correct translations are generated in a series of interactions between the ITP system and its user. Specifically, at each interaction of the ITP process, the ITP system generates a translation of the source sentence which can be partially or completely accepted and corrected by the user of the ITP system. Each partially corrected text segment (referred to from now on as prefix), is then used by the SMT system as additional information to generate better translation suggestions.

An example of a typical ITP session is shown in Figure 1.1. In interaction-0, the system suggests a translation (s). In interaction-1, the user moves the mouse to accept the prefix composed of the first eight characters “To view ” (p) and presses the **[a]** key (k), then the system suggests completing the sentence with “list of resources” (a new s). Interactions 2 and 3 are similar. In the final interaction, the user completely accepts the current suggestion.

Figure 1.1: ITP session to translate a Spanish sentence into English.

source (f_1^J):		Para ver la lista de recursos				
reference (\hat{e}_1^I):		To view a listing of resources				
interaction-0	p s	To view the resources list				
interaction-1	p k s	To view [a] list of resources				
interaction-2	p k s	To view a list [i] ng resources				
interaction-3	p k s	To view a listing [o] f resources				
acceptance	p	To view a listing of resources				

Figure 1.2 shows a schematic view of these ideas. Here, f_1^J is the input sentence and e_1^I is the output derived by the ITP system from f_1^J . By observing f_1^J and e_1^I , the user interacts with the ITP system, validating prefixes and/or pressing keys (k) corresponding to the next correct character, until the desired output \hat{e}_1^I is produced. The models used by the ITP system are

obtained through a classical batch training process from a previously given training sequence of pairs (f_n, e_n) from the task being considered.

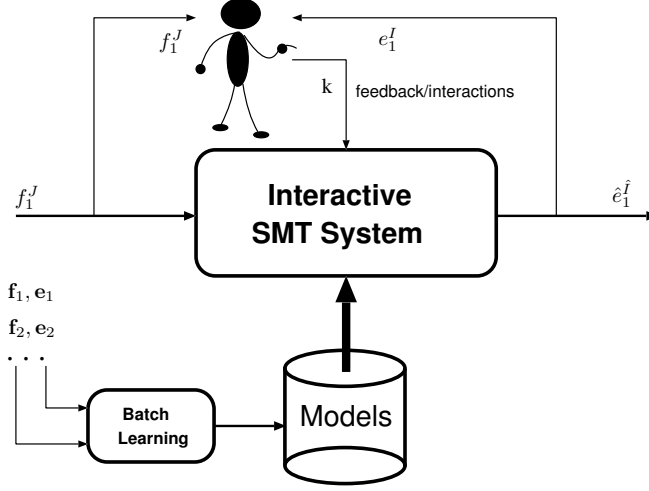


Figure 1.2: An Interactive SMT system.

More formally, in the ITP scenario we have to find an extension s for a prefix p given by the user:

$$\hat{s} = \arg \max_s \{p(s \mid f_1^J, p)\} \quad (1.4)$$

Applying the Bayes rule, we arrive at the following expression:

$$\hat{s} = \arg \max_s \{p(s \mid p) \cdot p(f_1^J \mid p, s)\} \quad (1.5)$$

where the term $p(p)$ has been dropped since it does not depend on s .

Thus, the search is restricted to those sentences e_1^I which contain p as prefix. It is also worth mentioning that the similarities between Equation (1.5) and Equation (1.2) (note that $ps \equiv e_1^I$) allow us to use the same models if the search procedures are adequately modified (Bender et al. 2005; Barrachina et al. 2009).

1.2 Toolkit Features

The toolkit includes the following features:

- Phrase-based statistical machine translation decoder.
- Computer-aided translation (post-edition and interactive machine translation).
- Incremental estimation of all of the models involved in the translation process.

- Client-server implementation of the translation functionality.
- Single word alignment model estimation using the incremental EM algorithm.
- Scalable and parallel model estimation algorithms using Map-Reduce.
- Compiles on Unix-like and Windows (using Cygwin) systems.
- Integration with the CasMaCat Workbench developed in the EU FP7 CasMaCat project^b.
- ...

1.3 Distribution Details

Thot has been coded using C, C++ and shell scripting. Thot is known to compile on Unix-like and Windows (using cygwin) systems. As future work we plan to port the code to other platforms. See Section 1.6 section of this file if you experience problems during compilation.

It is released under the GNU Lesser General Public License (LGPL)^c.

1.4 Relation with Existing Software

Due to the strong focus of Thot on online and incremental learning, it includes its own programs to carry out language and translation model estimation. Specifically, Thot includes tools to work with n -gram language models based on incrementally updateable sufficient statistics. On the other hand, Thot also includes a set of tools and a whole software library to estimate IBM 1, IBM 2 and HMM-based word alignment models. The estimation process can be carried out using batch and incremental EM algorithms. This functionality is not based on the standard GIZA++ software for word alignment model generation.

Additionally, Thot does not use any code from other existing translation tools such as Moses. In this regard, Thot is a smaller project which tries to offer its own view of the process of statistical machine translation, with a strong focus on online learning and also incorporating interactive machine translation functionality.

1.5 Current Status

The Thot toolkit is under development. Original public versions of Thot date back to 2005 (Ortiz et al. 2005) and did only include estimation of phrase-based models. By contrast, current version offers several new features that had not been previously incorporated.

A basic usage manual is currently being developed. In addition to this, a set specific tools to ease the process of making SMT experiments has been created.

In addition to the basic usage manual, there are some toolkit extensions that will be incorporated in the next few months:

^b<http://www.casmacat.eu/>

^c<http://www.gnu.org/copyleft/lgpl.html>

- Virtualized language models (i.e. accessing language model parameters from disk).
- Interpolation of language and translation models.
- Improved concurrency in the Thot translation server (translation process is not still concurrent).

Finally, here is a list of known issues with the Thot toolkit that are currently being addressed:

- Phrase model training is based on HMM-based alignments models estimated by means of incremental EM. This estimation process is computationally demanding and currently constitutes a bottleneck when training phrase models from large corpora. One already implemented solution is to carry out the estimation in multiple processors. However, we are also investigating alternative optimization techniques that allow us to efficiently execute the estimation process in a single processor.
- Log-linear model weight adjustment is carried out by means of the downhill simplex algorithm, which is very slow. Downhill simplex will be replaced by a more efficient technique.
- Non-monotonic translation is not yet sufficiently tested, specially with complex corpora such as Europarl.

1.6 Documentation and Support

Project documentation is being developed. Such documentation include:

- README file included with the Thot package.
- The Thot manual (`thot_manual.pdf` under the `doc` directory).
- Thot website^d.

If you need additional help, you can:

- use the github issue tracker^e.
- send an e-mail to the author^f.
- join the CasMaCat support group^g.

Additional information about the theoretical foundations of Thot can be found in:

^d<http://daormar.github.io/thot/>

^e<https://github.com/daormar/thot/issues>

^fdortiz@prhlt.upv.es

^g<http://groups.google.com/group/casmacat-support/boxsubscribe>

- Daniel Ortiz-Martínez. *Advances in Fully-Automatic and Interactive Phrase-Based Statistical Machine Translation*. *PhD Thesis*. Universitat Politècnica de València. Advisors: Ismael García Varea and Francisco Casacuberta. 2011.

One interesting feature of Thot incremental (or online) estimation of statistical models, is also described in the following paper:

- Daniel Ortiz-Martínez, Ismael García-Varea, Francisco Casacuberta. *Online learning for interactive statistical machine translation*. In Proc. of the North American Chapter of the Association for Computational Linguistics - Human Language Technologies (NAACL-HLT), pp. 546–554, Los Angeles, US, June 2010.

The phrase-level alignment generation functionality is described in:

- Daniel Ortiz-Martínez, Ismael García-Varea, Francisco Casacuberta. *Phrase-level alignment generation using a smoothed loglinear phrase-based statistical alignment model*. In Proc. of the European Association for Machine Translation (EAMT), pp. 160-169, Hamburg, Germany, 2008. *Best paper award*.

Finally, the initial version of Thot was described in:

- Daniel Ortiz-Martínez, Ismael García-Varea, Francisco Casacuberta. *Thot: a toolkit to train phrase-based models for statistical machine translation*. In Proc. of the Tenth Machine Translation Summit (MT-Summit), Phuket, Thailand, September 2005.

1.7 Citation

You are welcome to use the code under the terms of the license for research or commercial purposes, however please acknowledge its use with a citation:

- Daniel Ortiz-Martínez, Francisco Casacuberta. *The New Thot Toolkit for Fully Automatic and Interactive Statistical Machine Translation*. In Proc. of the 14th Annual Meeting of the European Association for Computational Linguistics (ACL): System Demonstrations, pp. 45–48, Gothenburg, Sweden, April 2014.

Here is a BiBTeX entry:

```
@InProceedings{Ortiz2014,
  author    = {D. Ortiz-Mart\'{\i}nez and F. Casacuberta},
  title     = {The New Thot Toolkit for Fully Automatic and
              Interactive Statistical Machine Translation},
  booktitle = {14th Annual Meeting of the European Association for Computational
              Linguistics: System Demonstrations},
  year      = {2014},
  month     = {April},
  address   = {Gothenburg, Sweden},
  pages     = "45--48",
}
```

1.8 Sponsors

Thot is currently supported by the European Union under the CasMaCat research project. Thot has also received support from the Spanish Government in a number of research projects, such as the MIPRCV project^h that belongs to the CONSOLIDER programmeⁱ.

^h<http://miprcv.iti.upv.es/>

ⁱ<http://www.ingenio2010.es/>

CHAPTER 2

INSTALLATION

2.1 Basic Installation Procedure

The code of the Thot toolkit is hosted on github^a. To install Thot, first you need to install the autotools (autoconf, autoconf-archive, automake and libtool packages in Ubuntu). If you are planning to use Thot on a Windows platform, you also need to install the Cygwin environment.

Once the autotools are available (as well as Cygwin if required), you can proceed with the installation of Thot by following the next sequence of steps:

1. Obtain the package using git:

```
$ git clone https://github.com/daormar/thot.git
```

Additionally, Thot can be downloaded in a zip file^b.

2. `cd` to the directory containing the package's source code and type `./reconf`.
3. Type `./configure` to configure the package
4. Type `make` to compile the package
5. Type `make install` to install the programs and any data files and documentation
6. You can remove the program binaries and object files from the source code directory by typing `make clean`

^a<https://github.com/>

^b<https://github.com/daormar/thot/archive/master.zip>

By default the files are installed under the `/usr/local` directory (or similar, depending of the OS you use); however, since Step 5 requires root privileges, another directory can be specified during Step 3 by typing:

```
$ configure --prefix=<absolute-installation-path>
```

For example, if `user1` wants to install the Thot package in the directory `/home/user1/thot`, the sequence of commands to execute should be the following:

```
$ ./reconf
$ configure --prefix=/home/user1/thot
$ make
$ make install
```

The installation process also creates three directories with additional information:

- **`${PREFIX}/share/thot/cfg.templates`**: contains configuration files to be used with different Thot utilities (see Chapter 3 for more details).
- **`${PREFIX}/share/thot/doc`**: contains the documentation of Thot, which currently consists in the Thot manual (`thot_manual.pdf`).
- **`${PREFIX}/share/thot/toy_corpus`**: contains a very small parallel corpus to make software tests. This corpus may also be useful for new Thot users trying to get familiar with the toolkit functionality.

See the `INSTALL` file in the directory where Thot has been downloaded for more information.

IMPORTANT NOTE: if Thot is being installed in a PBS cluster (a cluster providing `qsub` and other related tools), it is important that the `configure` script is executed in the main cluster node, so as to properly detect the cluster configuration (do not execute it in an interactive session).

2.2 Alternative Installation Options

The Thot configure script can be used to modify the toolkit behavior. Here is a list of current installation options:

- **`--enable-ibm2-align`**: Thot currently uses HMM-based alignment models to obtain the word alignment matrices required for phrase model estimation. One alternative installation option allows to replace HMM-based alignment models by IBM 2 alignment models. IBM 2 alignment models can be estimated very efficiently without significantly affecting translation quality.
- **`--with-casmacat`**: this options enables the configuration required for the CasMaCat Workbench, see more information below.

2.3 Installation Including the CasMaCat Workbench

Thot can be combined with the CasMaCatWorkbench which has been developed in the project of the same name. The specific installation instructions can be obtained at the project website^c.

^c<http://www.casmacat.eu/index.php?n=Workbench.Workbench>

CHAPTER 3

USER GUIDE

This chapter provides usage information for the Thot toolkit. A toolkit overview is given in Section 3.1. Model training and tuning tools are presented in Section 3.2. Section 3.3 explains how to use the previously trained and tuned models to generate translations or phrase alignments. Section 3.4 mentions some additional tools that are useful for translation tasks and Section 3.5 describes some advanced features implemented by Thot. Section 3.6 provides general sample uses of Thot. Finally, Section 3.7 gives troubleshooting information about the toolkit usage.

3.1 Toolkit Overview

The basic usage of Thot involves training, tuning and search processes. The training process consists in estimating the parameters of the translation and language models. After that, a basic configuration file collecting the data of the trained models is generated (the workflow implemented by the Thot toolkit makes extensive use of configuration files to increase usability). Once these models have been generated, they are combined by means of the so-called log-linear models. This combination assigns different weights to the models so as to increase the translation quality. The exact values of such weights are determined during the tuning process. After training and tuning the models, an optional filtering step is carried out so as to keep the portion of the phrase model that is strictly necessary to work with a specific test corpus. This filtering process may be crucial to ensure the applicability of statistical machine translation in real scenarios, due to the huge size of the phrase models that are obtained when processing large training corpora. Finally, during the search process, the resulting model is applied to generate translations in a fully-automatic or interactive way, or to generate alignments at phrase level. Such processes can be summarized in the following list:

1. Language model training.
2. Translation model training.
3. Generate basic configuration file.
4. Parameter tuning.
5. Phrase model filtering (optional).
6. Search:
 - (a) Fully automatic translation.
 - (b) Interactive machine translation.

- (c) Phrase alignment generation.

Thot allows us to execute in parallel the above explained tasks using computer clusters or multi-processor systems. Parallel implementation is transparent to the user, who is requested to specify the number of processors in which the tools will be executed. Thot currently supports the use of PBS clusters (a cluster providing `qsub` and other related tools).

In the following sections we describe the different tools offered by the Thot toolkit to fully exploit its functionality.

3.2 Training and Tuning Tools

In this section, the functionality implemented in Thot for model training and tuning is described.

3.2.1 Language Model Training

Thot uses n -gram models with Jelinek-Mercer smoothing (see for instance (Chen and Goodman 1996)) to implement language models^a.

Basic Tools

Thot provides the `thot_lm_train` tool, which can be used to train language models. The basic input parameters include:

- pr** `<int>` : number of processors used to perform the estimation.
- c** `<string>` : monolingual corpus used to estimate the model.
- o** `<string>` : directory for output files. When executing the command in PBS clusters, the user should ensure that the provided path is accessible for all nodes involved in the computation.
- n** `<int>` : order of the n -grams.

The command also generates a model descriptor in the output directory name `lm_desc` that will be useful to generate configuration files.

When executing the estimation process in computer clusters, it is important to ensure that the computer nodes involved receive enough resources (memory, computation time, etc.). For this purpose, the `thot_lm_train` tool incorporates the `-qs` option^b. In addition to this, `-tdir` and `-sdir` options allow to set the paths of the directory for temporary and shared files, respectively, when the default ones do not have enough free space to carry out the estimation.

For additional options and information, command help can be obtained by typing `thot_lm_train --help`.

Examples

To illustrate the tools related to language model generation, as well as for other examples shown in this chapter, we will use the toy corpus included with the Thot toolkit. Assuming that Thot was installed in `${PREFIX}` directory, the toy corpus will be available in `${PREFIX}/share/thot/toy-corpus`^b.

^aState-of-the-art Kneser-Ney smoothing is currently being incorporated.

^bIf `--prefix` option was not explicitly provided to `configure` during installation, then `${PREFIX}` is set to `/usr/local`.

The following command line trains a 3-gram language model for the Spanish training set of the toy corpus, using the `-unk` option, storing the results in the `lm_outdir` directory:

```
train_corpus=${PREFIX}/share/thot/toy_corpus/train.sp
thot_lm_train -c ${train_corpus} -o lm_outdir -n 3 -unk
```

3.2.2 Translation Model Training

Thot implements translation models by means of phrase-based models (Koehn et al. 2003).

Basic Tools

Thot incorporates the `thot_tm_train` tool, useful to train phrase models. The basic input parameters include:

- pr** `<int>` : number of processors used to perform the estimation.
- s** `<string>` : file with source sentences.
- t** `<string>` : file with target sentences.
- o** `<string>` : directory for output files. When executing the tool in PBS clusters, the user should ensure that the provided path is accessible for all nodes involved in the computation.

The command also generates a model descriptor in the output directory name `tm_desc` that can be used to generate configuration files.

`thot_tm_train` also includes options to ensure that the training process receives enough computational resources in the same way as was explained for language model training (see Section 3.2.1). These options are `-qs`, `-tdir` and `-sdir`.

For additional options and information, command help can be obtained by typing `thot_tm_train --help`.

Examples

Again, we will use the toy corpus included with the Thot toolkit to illustrate the translation model functionality. Assuming that Thot was installed in `${PREFIX}` directory, the toy corpus will be available in `${PREFIX}/share/thot/toy_corpus`.

The following command line trains a phrase model for the Spanish training set of the toy corpus, storing the results in the `tm_outdir` directory:

```
src_train_corpus=${PREFIX}/share/thot/toy_corpus/train.sp
trg_train_corpus=${PREFIX}/share/thot/toy_corpus/train.en
thot_tm_train -s ${src_train_corpus} -t ${trg_train_corpus} -o tm_outdir
```

3.2.3 Basic Configuration File Generation

Thot uses configuration files to simplify the access to the toolkit functionality. These configuration files provides information about the parameters used by the Thot translation system, including the location of the translation and language models, the set of log-linear model weights, parameters relevant to the way in which the search process is carried out, etc.

Configuration files can be manually generated from the template files given in the path `${PREFIX}/share/thot/cfg-templates`, where `${PREFIX}` is the directory where Thot was installed.

Basic Tools

Alternatively, configuration files can also be generated by means of the `thot_gen_cfg_file` command. For this purpose, it is necessary to provide the descriptors of both the language and translation models (the exact syntax can be obtained by executing the command without parameters). The command output is written to the standard output and consists in a basic configuration file allowing to work with the language and translation models given by the provided model descriptors.

Examples

Assuming that we have already trained language and translation models located in the `lm_outdir` and `tm_outdir` directories, respectively, the following command line generates a basic Thot configuration file and writes it to the file `example.cfg`:

```
thot_gen_cfg_file lm_outdir/lm_desc tm_outdir/tm_desc > example.cfg
```

3.2.4 Parameter Tuning

After training the language and translation models, it is necessary to execute a parameter tuning stage. For this purpose, a development corpus separated from the training corpus is required. Currently, this stage affects to the weights of the language model as well as those of the log-linear model. Tuning of language model weights is necessary due to the use of Jelinek-Mercer smoothing (as it was mentioned above, Kneser-Ney smoothing is being implemented).

Thot incorporates the downhill-simplex algorithm (Nelder and Mead 1965) to tune the language and log-linear model weights. Regarding the criterion used for weight adjustment, language models weights are set so as to minimize the perplexity of the model, while the criterion to adjust log-linear weights is to maximize translation quality in terms of the well known BLEU measure (Papineni et al. 2001).

Basic Tools

The `thot_smt_tune` tool allows to perform parameter tuning. For this purpose, it is necessary a Thot configuration file and a development corpus. Here is a list of the basic input parameters:

- `-pr` `<int>` : number of processors used to perform the estimation.
- `-c` `<string>` : Thot configuration file.
- `-s` `<string>` : file with source sentences.

-t **<string>** : file with target sentences.

-o **<string>** : directory for output files. When executing the command in PBS clusters, the user should ensure that the provided path is accessible for all nodes involved in the computation.

`thot_smt_tune` returns a new configuration file where the language and log-linear model weights are tuned. This file is stored under the output directory given by the `-o` option with the name `tuned_for_dev.cfg`.

As in previously presented tools, `thot_smt_tune` may require specific resources that can be specified by means of the `-qs`, `-tdir` and `-sdir` options.

For additional options and information, command help can be obtained by typing `thot_smt_tune --help`.

Examples

The following command line tunes the system given in the `example.cfg` file, for the development set of the Thot toy corpus, storing the results in the `tune` directory:

```
src_dev_corpus=${PREFIX}/share/thot/toy_corpus/dev.sp
trg_dev_corpus=${PREFIX}/share/thot/toy_corpus/dev.en
thot_smt_tune -c example.cfg -s ${src_dev_corpus} -t ${trg_dev_corpus} \
-o tune
```

After the successful execution of `thot_smt_tune`, the configuration file corresponding to the tuned system will be stored in `tune/tuned_for_dev.cfg`.

3.2.5 Phrase Model Filtering

Phrase models are composed of millions of parameters when they are estimated from large training corpora, making impossible to store them in main memory when using regular computer hardware. One simple solution to this problem when the set of sentences to be translated is known beforehand is to filter those phrase model parameters that are relevant to carry out the translation process.

Basic Tools

The `thot_prepare_sys_for_test` tool allows to filter the parameters of a phrase model. For this purpose, it is necessary a Thot configuration file and a file with the sentences to be translated. Here is a list of the basic input parameters:

-c **<string>** : Thot configuration file.

-t **<string>** : file with test sentences.

-o **<string>** : directory for output files. When executing the command in PBS clusters, the user should ensure that the provided path is accessible for all nodes involved in the computation.

The `thot_prepare_sys_for_test` tool may require specific resources that can be specified by means of the `-qs`, `-tdir` and `-sdir` options, in a similar way to other tools described above.

For additional options and information, command help can be obtained by typing `thot_prepare_sys_for_test --help`.

Examples

The following command line filter the phrase model given in the `example.cfg` file, for the test set of the Thot toy corpus, storing the results in the `filtdir` directory:

```
src_test_corpus=${PREFIX}/share/thot/toy_corpus/test.sp
thot_prepare_sys_for_test -c example.cfg -t ${src_test_corpus} -o filtdir
```

3.3 Search Tools

After performing model training, tuning and filtering, we are prepared to generate translations, both in a fully-automatic or interactive way. In addition to this, the estimated models can also be used to generate phrase level alignments. The following sections explain how to access such functionalities.

3.3.1 Fully Automatic Translation

Thot incorporates tools to translate a set of sentences in a fully-automatic way. This task can be carried out using both, command line and client-server tools.

Basic Tools

The `thot_decoder` tool allows to generate translations for a given test set. For this purpose, a Thot configuration file and a file with the sentences to be translated should be provided. Here is a list of the basic input parameters:

`-c` `<string>` : Thot configuration file.

`-t` `<string>` : file with test sentences.

`-o` `<string>` : output file.

For additional options and information, command help can be obtained by typing `thot_decoder --help`.

Client-Server Tools

The translation functionality mentioned above is also included in the client-server architecture provided by Thot. This includes two basic tools: `thot_server` and `thot_client`.

The `thot_server` tool implements a fully-fledged SMT system. The most relevant input parameter that has to be provided, the `-c` parameter, is the name of the configuration file.

On the other hand, `thot_client` can be used to request translations to the server. `thot_client` requires the IP address where the server is being executed using the `-i` option, as well as the sentence to be translated, that is provided by means of the `-t` option.

Both the `thot_client` and `thot_server` tools may receive additional input parameters, use the `--help` option for more information.

Examples

The following command uses the system configuration provided in the `example.cfg` file to translate the test set of the Thot toy corpus, storing the results in the `filtdir` directory:

```
src_test_corpus=${PREFIX}/share/thot/toy_corpus/test.sp
thot_decoder -c example.cfg -t ${src_test_corpus} -o filtdir
```

Alternatively, the following example shows how to obtain the translation of a test sentence of the toy corpus using the client-server architecture:

```
thot_server -c example.cfg & # start server
thot_client -i 127.0.0.1 -t "voy a marcharme hoy por la tarde ."
```

3.3.2 Interactive Machine Translation

Thot implements interactive machine translation functionality, which allows to generate suffixes that complete the prefixes validated by the user (see Section 1.1.2).

Client-Server Tools

The interactive machine translation functionality implemented in Thot can be accessed by means of the `thot_server` and `thot_client` tools. This functionality includes obtaining the initial translation of a source sentence and obtaining the suffix that best completes a prefix given by the user.

Before using the interactive machine translation functionality, it is necessary to initialize the server. For this purpose, the `thot_server` tool requires the name of configuration file that is supplied by means of the `-c` parameter.

On the other hand, requests to the server can be sent by means of `thot_client`. As basic parameter, `thot_client` requires the IP address where the server is being executed using the `-i` option. Thot maintains state information through the different interactions between the user and the system. Some of the information that is maintained is specific to the user and hence it is important to use different user identifiers for different interactive translation sessions. The user identifier can be provided by means of the `-uid` option. To interactively translate a sentence, the `-sc` option is used to start the translation process. After that, new strings can be added to the previously existing prefix by means of the `-ap` option. It is important to point out that the server works at character level. Finally, the `-rp` option can be used to reset the prefix.

To get additional information on the usage of `thot_client` and `thot_server`, use the `--help` option.

Examples

The following example shows how to access the basic interactive machine translation functionality provided by Thot using the client-server architecture:

```
thot_server -c example.cfg & # start server
thot_client -i 127.0.0.1 -uid 0 -sc "me marchó hoy por la tarde ."
thot_client -i 127.0.0.1 -uid 0 -ap "I am"
thot_client -i 127.0.0.1 -uid 0 -ap " le"
# NOTE: at this point, the user prefix provided to the system is "I am le"
thot_client -i 127.0.0.1 -uid 0 -rp # reset prefix
```

3.3.3 Phrase Alignment Generation

TO-BE-DONE.

3.4 Additional Tools

In the following sections we describe some tools relevant to the translation process that were not listed above.

3.4.1 Output Evaluation

After translating a test set, the translation quality can be evaluated using automatic measures provided that there exist reference translations for each source sentence. Thot implements some tools for this purpose:

- **thot.calc.bleu**: obtains the BLEU (bilingual evaluation understudy) measure Papineni et al. 2001.
- **thot.calc.wer**: calculates the WER (word error rate) measure (the number of substitutions, insertions and deletions that are required to convert the system translation into the reference sentence).

3.5 Advanced Functionality

This section explains how to access advanced features included in the Thot toolkit.

3.5.1 Online Learning

Thot incorporates techniques that allow to incrementally update the parameters associated to the features of a state-of-the-art log-linear model (Ortiz-Martínez et al. 2010). For this purpose, a set of incrementally updateable sufficient statistics is defined for each feature, allowing us to process individual training samples in constant time complexity.

Client-Server Tools

The online learning functionality implemented in Thot can be accessed by means of the of the `thot_server` and `thot_client` tools. This functionality includes processing a single training pair and printing the models.

Prior to use the online learning functionality, the server should be initialized. For this purpose, the `thot_server` tool requires the name of configuration file that is supplied by means of the `-c` parameter.

On the other hand, requests to the server can be sent by means of `thot_client`. As basic parameter, `thot_client` requires the IP address where the server is being executed using the `-i` option. In addition to this, the `-tr` option can be used to specify the training pair to be processed and the `-pr` option allows to print the updated models to files. Printing the models causes the previous ones to be overwritten.

To get more information on the usage of `thot_client` and `thot_server`, use the `--help` option.

Examples

The following example shows how to process a new training pair and print the models using the client-server architecture:

```
thot_server -c example.cfg & # start server
thot_client -i 127.0.0.1 -tr "esto es una prueba" "this is a test"
thot_client -i 127.0.0.1 -pr # print models (previous ones are overwritten)
```

3.6 General Sample Uses

In this section we will show some general sample uses illustrating the functionality of the Thot toolkit.

3.6.1 Training, Tuning and Translating

Example 1: processing the toy corpus

This example shows the sequence of commands required to train, tune and translate using the toy corpus included with Thot.

```
# define variables (optional)
src_train_corpus=${PREFIX}/share/thot/toy_corpus/train.sp
trg_train_corpus=${PREFIX}/share/thot/toy_corpus/train.en
src_dev_corpus=${PREFIX}/share/thot/toy_corpus/dev.sp
trg_dev_corpus=${PREFIX}/share/thot/toy_corpus/dev.en
src_test_corpus=${PREFIX}/share/thot/toy_corpus/test.sp
trg_test_corpus=${PREFIX}/share/thot/toy_corpus/test.en

# train system
thot_lm_train -c ${trg_train_corpus} -o lm_outdir -n 3 -unk
thot_tm_train -s ${src_train_corpus} -t ${trg_train_corpus} -o tm_outdir

# generate cfg file
thot_gen_cfg_file lm_outdir/lm_desc tm_outdir/tm_desc > example.cfg
```

```
# tune system
thot_smt_tune -c example.cfg -s ${src_dev_corpus} -t ${trg_dev_corpus} \
-o tune

# filter phrase model
thot_prepare_sys_for_test -c tune/tuned_for_dev.cfg -t ${src_test_corpus} \
-o filtdir

# translate test corpus
thot_decoder -c filtdir/test_specific.cfg -t ${src_test_corpus} -o output

# evaluate translation quality
thot_calc_bleu -r ${trg_test_corpus} -t output
```

3.6.2 Online Learning

Example 1: adding a training pair to the toy corpus

The following example shows the commands that have to be executed to add a new training pair to the toy corpus. For this purpose, models for such corpus are first trained, tuned and filtered. After that, the Thot server is started and the client is used to incorporate the new training sample. Finally, the client is used again to print the resulting models to files, overwriting the previous ones.

```
# define variables (optional)
src_train_corpus=${PREFIX}/share/thot/toy_corpus/train.sp
trg_train_corpus=${PREFIX}/share/thot/toy_corpus/train.en
src_dev_corpus=${PREFIX}/share/thot/toy_corpus/dev.sp
trg_dev_corpus=${PREFIX}/share/thot/toy_corpus/dev.en
src_test_corpus=${PREFIX}/share/thot/toy_corpus/test.sp
trg_test_corpus=${PREFIX}/share/thot/toy_corpus/test.en

# train system
thot_lm_train -c ${trg_train_corpus} -o lm_outdir -n 3 -unk
thot_tm_train -s ${src_train_corpus} -t ${trg_train_corpus} -o tm_outdir

# generate cfg file
thot_gen_cfg_file lm_outdir/lm_desc tm_outdir/tm_desc > example.cfg

# tune system
thot_smt_tune -c example.cfg -s ${src_dev_corpus} -t ${trg_dev_corpus} \
-o tune

# filter phrase model
thot_prepare_sys_for_test -c tune/tuned_for_dev.cfg -t ${src_test_corpus} \
-o filtdir

# start Thot server
thot_server -c filtdir/test_specific.cfg &

# add new training pair
```

```
thot_client -i 127.0.0.1 -tr "esto es una prueba" "this is a test"

# print models (warning: previously generated models are overwritten)
thot_client -i 127.0.0.1 -pr
```

3.7 Troubleshooting

This section provides troubleshooting information about possible problems that arise during the toolkit usage. The current list of identified problems is the following:

- **The `thot_tm_train` tool is too slow:** Thot uses HMM-based alignment models to obtain the word alignment matrices required for phrase model estimation. The current implementation of this kind of models is slow and may constitute a bottleneck. To alleviate this problem, toolkit users may enable one of the following workarounds:
 - a) Use the `-pr` option to execute `thot_tm_train` in multiple processors.
 - b) Replace HMM-based alignment models by IBM 2 alignment models by means of the `--enable-ibm2-align` option of the `configure` script (see Section 2.2 for more details). IBM 2 alignment models can be estimated and subsequently used to generate word alignment matrices very efficiently without causing significant degradations in the translation quality.

CHAPTER 4

BACKGROUND

TO-BE-DONE

BIBLIOGRAPHY

- Barrachina, S. et al. (2009). “Statistical Approaches to Computer-Assisted Translation”. In: *Computational Linguistics* 35.1, pp. 3–28.
- Bender, O., S. Hasan, D. Vilar, R. Zens, and H. Ney (2005). “Comparison of Generation Strategies for Interactive Machine Translation”. In: *Conference of the European Association for Machine Translation*. Budapest, Hungary, pp. 33–40.
- Brown, P. F., S. A. Della Pietra, V. J. Della Pietra, and R. L. Mercer (1993). “The Mathematics of Statistical Machine Translation: Parameter Estimation”. In: *Computational Linguistics* 19.2, pp. 263–311.
- Chen, S. F. and J. Goodman (1996). “An Empirical Study of Smoothing Techniques for Language Modeling”. In: *Proceedings of the 34th Annual Meeting of the Association for Computational Linguistics*. Ed. by Arivind Joshi and Martha Palmer. San Francisco: Morgan Kaufmann Publishers, pp. 310–318.
- Garcia, I. (2011). “Translating by post-editing: is it the way forward?” In: *Machine Translation* 25.3, pp. 217–237. ISSN: 0922-6567.
- Koehn, P., F. J. Och, and D. Marcu (2003). “Statistical Phrase-Based Translation”. In: *Proceedings of the Human Language Technology and North American Association for Computational Linguistics Conference*. Edmonton, Canada, pp. 48–54.
- Nelder, J. A. and R. Mead (1965). “A Simplex Method for Function Minimization”. In: *The Computer Journal* 7.4, pp. 308–313. DOI: [10.1093/comjnl/7.4.308](https://doi.org/10.1093/comjnl/7.4.308).
- Ney, H. (1995). “On the Probabilistic-Interpretation of Neural-Network Classifiers and Discriminative Training Criteria”. In: *IEEE Transactions on Pattern Analysis and Machine Intelligence* 17.2, pp. 107–119.
- Och, Franz Josef (2003). “Minimum error rate training in statistical machine translation”. In: *Proceedings of the 41th Annual Conference of the Associations for Computational Linguistics*. Sapporo, Japan, pp. 160–167.
- Ortiz, D., I. García-Varea, and F. Casacuberta (2005). “Thot: a Toolkit To Train Phrase-based Statistical Translation Models”. In: *Proceedings of the Machine Translation Summit X*. Phuket, Thailand, pp. 141–148.
- Ortiz-Martínez, D., I. García-Varea, and F. Casacuberta (2010). “Online Learning for Interactive Statistical Machine Translation”. In: *Proceedings of the North American Chapter of the Association for Computational Linguistics - Human Language Technologies (NAACL-HLT)*. Los Angeles, pp. 546–554.
- Papineni, K. A., S. Roukos, T. Ward, and W. Zhu (2001). *Bleu: a Method for Automatic Evaluation of Machine Translation*. Tech. rep. RC22176 (W0109-022). Yorktown Heights, NY: IBM Research Division, Thomas J. Watson Research Center, 10 pages.
- TAUS-Project (2010). *Postediting in Practice. A TAUS Report*. Tech. rep. TAUS – Enabling better translation. URL: <https://www.taus.net/reports/postediting-in-practice>.